Technology

When Chatbots Choose

Who does AI really serve?

Alex Sergeev

In September, I needed an alarm clock. Considering the very large, highly competitive market with numerous similar listings, I decided to delegate this task to ChatGPT. It found a clock that matched my constraints. But when AI chooses, who is it really working for?

We can view this through the lens of the Principal–Agent Problem. The consumer delegates search and selection to an agent, without easily being able to observe the agent's incentives. Because the AI is the broker between buyers and sellers, it becomes a gatekeeper, much like existing online marketplaces such as Amazon, which monetise attention through advertising.

For consumers

For a consumer such as myself, the incentives are clear. Following Hull's "Law of Least Effort", we want to outsource the research. We want to minimise time spent and attention while getting a "good enough" product. Considering it is very difficult to get the best product possible — sometimes requiring a couple of hours of searching — for a £20 item such as an alarm clock, it would be rational to use ChatGPT to make the choice instead, as the benefit is very small compared to the effort.

Consumers only benefit if AI really is optimising the constraints. In the case of the alarm clock, this would be noise, accuracy, and price. However, consumers cannot verify that the chatbot is actually doing this. As AI systems are largely "black boxes" and prone to "hallucinations", the user cannot easily see the ranking criteria or the extent of the search. To verify the recommendation, consumers must do the very search they outsourced. Furthermore, firms could develop methods of "cheating" to influence rankings, like how they used invisible words to manipulate older search engines. Firms will invest in features that the AI can process and reward, sometimes prioritising this over true product quality as the market responds to them. It becomes a race to be one of the products the AI recommends — such as one of the three alarm clocks I was given — rather than a race to be the best in actual quality. Fake reviews would further reduce consumer welfare while increasing perceived quality within rankings.

What about AI?

The service itself wants revenue, and depending on the funding model, it could pursue different monetisation strategies. If it operates under a subscription-based model, incentives are more aligned with the consumer: the service (e.g. OpenAI) wants to retain paying users and therefore has reason to deliver recommendations that genuinely meet user needs. The conflict arises when services monetise through advertising. In this case, there may be a significant conflict of interest. The service may be paid by firms seeking to sell more products and could prioritise these over others for that reason alone. We already see this with Amazon, where sponsored results appear above organic listings and divert attention.

This issue may be even more pronounced in AI services, where recommendations are limited to only a handful of products, potentially excluding non-sponsored options entirely and creating barriers for sellers. Firms may be forced to pay for visibility, costs that are likely to be passed on to consumers. This would significantly reduce consumer welfare through both higher prices and reduced choice.

For the Regulator

If AI becomes the primary gatekeeper allocating consumer attention, regulation will become increasingly important for protecting consumer welfare. To reduce information asymmetry, policies supporting transparency and independent audits would be effective, alongside providing lists of sources. This would allow regulators to examine how products are analysed, even if only through sampled evaluations.

To address conflicts of interest in the principal–agent problem, a simple rule — already applied by platforms such as Amazon — is the clear labelling of “sponsored” products. In practice, this would allow consumers to identify whether AI recommendations are influenced by payment or bias. As for the gatekeeper issue, existing antitrust rules and limits on self-preferencing could be applied. This would enable consumers to assess which assistants are most trustworthy based on public opinion and disclosed processes and make informed choices accordingly.

Conclusion

We delegate our choices with increasing frequency. Outsourcing decisions lowers search costs but shifts power to recommender AI systems and their owners. Currently, as services focus on growth, there are gains in consumer welfare. However, in the future, competition for visibility and advertising incentives may distort markets, raise prices, and reduce consumer welfare. To regulate this, incentive disclosure and independent auditing will be essential. Regulators, it seems, may need their own alarm clock.